

Bruno Rodríguez

Data infrastructure

+1 (628) 502-7009 | bruno.rodriguez.sep@gmail.com | brunorodriguez.dev | [in](#) [bruno-rodriguez-sepulveda](#) | H-1B1 visa

Experience

Breezy, property data for real estate

San Francisco, USA

Founding Engineer (2nd)

04/2025 – Present

- Owned zoning extraction for **11,569 US cities** in OpenTofu: Step Functions, Lambda, Fargate, and SQS, with Aurora, DynamoDB, and Bedrock behind least-privilege IAM and SNS alarms. Landed 15,113 zone records
- Moved parcel-to-zoning assignment onto a dedicated ECS task over ClickHouse (H3 candidate index, exact point-in-polygon, DynamoDB lease). Cut assignment from **33–41 min to ~11 min**; national rollups finish in 31–38 min, with the timeout budget promoted dev → staging → prod
- Built nationwide place search: Glue snapshots into OpenSearch, shadow-index publish with fail-closed abort and N-1 rollback, and reader IAM for the query service. Serving benchmark **9 ms p95**. Relevance checked on a 327-case frozen holdout (~2,200 runs)
- Operated coverage snapshots on S3 and DynamoDB, refreshed by a constrained `ecs:RunTask`, with Casbin RBAC and audit logs across three services. Datadog alarms on snapshot age, artifact size, and task memory
- Imported Auth0 dev and prod into OpenTofu and replaced main-branch prod deploys with a tag-triggered pipeline (GitHub OIDC, dry-run by default, concurrency lock against the legacy workflow)

Fintual (YC S18, Sequoia Capital)

Remote, Worldwide

Software Engineer, Onboarding & Recurrence

04/2024 – 04/2025

- Built an automated tax-verification pipeline that reduced manual processing **80%**, from ~50 to ~10 minutes per day
- Rebuilt onboarding and ran an A/B test that increased conversion **6 percentage points** across a 150K-user product
- Put typed TypeScript contracts on the legacy React Native API surface

Millennium Institute for Foundational Research on Data

Santiago, Chile

ML Infrastructure & Privacy Researcher

12/2022 – 03/2024

- Built a Spark/MLlib pipeline over ~1B records and the evaluation tooling around it, cutting model-analysis time **30%**
- Designed parallel, privacy-preserving clustering on that pipeline, improving precision **21%** over the baseline

Isatec, R&D

Santiago, Chile

Software Engineer & Team Lead (Part-time)

07/2022 – 08/2023

- Led eight engineers on a production exam pipeline; cut mobile-to-server sync from **10 minutes to 1 minute** per patient

Skills

Code	TypeScript, Python
Infrastructure	AWS (Step Functions, Lambda, ECS/Fargate, Glue, Aurora, DynamoDB, S3, CloudFront, IAM), OpenTofu, GitHub Actions
Data systems	ClickHouse, OpenSearch, PostgreSQL, H3 · LLM batch extraction, Spark/MLlib
Languages	English (C1), French (C1), Spanish (native)

Education

University of Chile

Santiago, Chile

BS + MEng in Computer Science · GPA 3.94 / 4.0 · Rank 10 / 116 · *Summa Cum Laude*

2017 – 2023

École CentraleSupélec, Paris-Saclay University

Paris, France

Diplôme d'ingénieur (MSc equivalent) · Industrial Engineering & Computer Science · Dual degree

2019 – 2022

- Honors:** Chilean national olympiads: 2× Gold in math teams, Silver and Bronze in math, Silver in physics; chemistry and informatics finalist. Selected for the Ibero-American Physics Olympiad
- Top 1% nationally on Chile's university admission exam; highest mathematics score in the Atacama Region
- Selected for Bayer Kimlu, an international science and leadership camp (3% acceptance)